

FIWB AI

Privacy Policy

Effective Date: February 2026 | Version 1.0

Operated by SIOWA | Owais Sayyed & Siddhant Wagh

Who This Policy Applies To

This Privacy Policy applies to all users of FIWB AI, a personalized academic intelligence platform. By connecting your institutional accounts or using the platform, you acknowledge that you have read and understood this policy. If you do not agree, please discontinue use.

1. What We Collect and Why

FIWB AI is built on the principle of contextual personalization. To deliver accurate, academically-relevant responses, we collect data from several sources. The table below summarizes each data category, its source, and its purpose.

Data Category	Source	Purpose
Course names, assignments, materials, due dates, announcements	Google Classroom / Moodle LMS (via OAuth 2.0 or API token)	Build your Academic Vault for context-aware tutoring
Email content (Primary inbox, last 14 days)	Gmail (via OAuth 2.0 with read scope)	Extract academic appointments, project context, and time-sensitive updates
Documents (PDFs, Google Docs, text files)	Google Drive (user-selected folders or root)	Index personal study material for RAG retrieval
Uploaded files (PDFs, images, text)	Direct in-chat upload	Provide immediate contextual assistance per conversation thread
Conversation history (last 10 messages)	FIWB AI chat sessions	Maintain coherent multi-turn dialogue
Learning style, knowledge gaps, strengths	Synthesized from your interactions over time	Build and continuously refine your Digital Twin profile
Google OAuth tokens (access + refresh)	Google Sign-In flow	Enable background sync of Classroom and Gmail
Moodle URL and API token	Provided manually by user	Access institutional LMS course content

Usage and cost statistics	System-generated	Internal analytics for performance monitoring and cost control
Aggregated academic interest signals (derived from prompts)	FIWB AI chat sessions	Serve contextually relevant in-platform advertisements (free tier only; opt-out available)

2. How We Use Your Data

2.1 Core Product Functionality

Every query you submit triggers a parallel retrieval phase across five context streams: your Academic Vault (synced course materials), Assistant Intel (Gmail-derived context), Chat Assets (uploaded files), Enhanced Memories (insights from past conversations), and your User Profile (your Digital Twin). These streams are synthesized to generate responses that are precise, personalized, and grounded in your actual academic environment.

2.2 Digital Twin Synthesis

Following every interaction, a background Synthesis Agent analyzes the conversation to extract signals about your learning style, identify knowledge gaps, and recognize areas of mastery. This data is formatted into a persistent User Profile — your Digital Twin — that evolves over time to make every future interaction more accurate and relevant.

2.3 Personalized Advertising (Freemium Model)

FIWB AI operates on a freemium basis. To sustain the free tier, we may use signals derived from your in-app interactions — such as the academic subjects you study, the types of questions you ask, and your general field of interest — to serve contextually relevant advertisements within the platform. This means ads may reflect your academic context (for example, a computer science student may see ads for developer tools or internship programs).

We want to be transparent about what this means in practice:

- Ad targeting is based on aggregated academic context derived from your prompts — not on raw personal data, email content, or data sold to external brokers.
- Your data is never shared with advertisers directly. Advertisers see only that their ad was shown to a relevant academic audience segment.
- You may opt out of personalized ads at any time via account settings. Opting out will result in generic, non-targeted ads being shown instead.
- Users on a paid (ad-free) subscription tier will not be shown any advertisements and their data will not be used for ad personalization.

Your Choice

Personalized ads are on by default for free-tier users. You can turn this off in Settings > Privacy > Personalized Ads at any time. This will not affect your access to any core FIWB AI features.

2.4 What We Do Not Do

- We do not sell, rent, or trade your raw personal data to any third party.
- We do not use your data to train general-purpose AI models.
- We do not access email outside your primary inbox or beyond the 14-day lookback window.
- We do not read emails that have not been classified as academically or assistively relevant.
- We do not share identifiable personal data with advertisers.

3. AI Models and Third-Party Processing

FIWB AI uses two OpenAI language models to process your queries and generate responses:

Model	Role	Data Exposed
GPT-4o-mini	Triage & memory synthesis	Raw user query and minimal metadata
GPT-4o	Primary response generation	Current query, retrieved context from all 5 streams, last 10 messages of conversation history

All processing via OpenAI is governed by OpenAI's API usage policies. Your data is processed to complete your request and is not used to train OpenAI's models under their API terms.

We also use Supermemory as our vector database provider. All documents stored in Supermemory are tagged with your unique user_id and are strictly isolated — no data is shared between users.

4. Data Storage and Security

4.1 Local Database (app.db — SQLite)

The following data is stored in a local SQLite database on the FIWB AI server:

- Your standardized email address (used as your unique identifier)
- Google OAuth access and refresh tokens, including token expiry timestamps
- Moodle instance URL and API token (if connected)
- Raw conversation history, including Base64-encoded attachment previews
- Course metadata: names, instructor details, assignment due dates
- Usage statistics: API call counts and per-request cost tracking in USD

4.2 Vector Database (Supermemory)

The following data is stored as vector embeddings in Supermemory for RAG retrieval:

- Full text of synced course materials and resources
- Email-derived academic insights (not raw email text)
- Uploaded file content
- Your synthesized Digital Twin profile (learning style, gaps, strengths)

Every document in Supermemory is tagged with your user_id, source type, data category, and timestamp. Retrieval queries are strictly filtered by user_id, ensuring zero cross-user data leakage.

4.3 Email Address Standardization

To prevent duplicate or ghost profiles, all user identities are normalized through a standardize_email process. For example, user+alias@gmail.com is mapped to user@gmail.com. This ensures all your data is consistently associated with a single identity.

4.4 Real-Time Sync

Gmail sync is maintained via Google Pub/Sub webhooks. This means our system is notified of new emails in real time to keep your academic context current. Only emails within the primary inbox and within the 14-day lookback window are processed.

5. Your Rights and Controls

You have the following rights with respect to your data:

- Right to Access: You may request a copy of all data we hold about you, including your Digital Twin profile and conversation history.
- Right to Deletion: You may request deletion of your account and all associated data, including your SQLite records and Supermemory embeddings. Deletion requests will be fulfilled within 30 days.
- Right to Revoke Access: You may revoke FIWB AI's Google OAuth access at any time via your Google Account settings (myaccount.google.com). This immediately disables Gmail and Classroom sync.
- Right to Disconnect LMS: You may remove your Moodle credentials from the platform at any time via account settings.
- Right to Opt Out of Memory Synthesis: Users may request that background Digital Twin synthesis be disabled. Note that this will significantly reduce response personalization.

How to Submit a Data Request

To exercise any of the above rights, contact us at the email address listed in Section 8 of this policy. Please include your registered email address and a description of your request. We will acknowledge your request within 7 days and fulfil it within 30 days.

6. Data Retention

We retain your data for as long as your account is active. Specific retention periods are as follows:

- Conversation history: Retained indefinitely to power longitudinal memory, unless you request deletion.
- Gmail-derived insights: Refreshed on a rolling 14-day basis. Raw email text is never stored.
- OAuth tokens: Retained until revoked or account deletion.
- Digital Twin profile: Retained and continuously updated until account deletion.
- Usage and cost statistics: Retained for 12 months for internal auditing.

7. Children and Student Privacy

FIWB AI is designed for use by college and university students, who are typically 18 years of age or older. We do not knowingly collect personal data from individuals under the age of 18. If you believe a minor has provided us with personal data, please contact us immediately for prompt removal.

For users accessing FIWB AI through an institutional arrangement, the institution's own data processing agreements may apply in addition to this policy. We are committed to working with institutions to meet applicable educational data protection requirements.

8. Policy Updates and Contact

We may update this Privacy Policy from time to time as the platform evolves. Where changes are material, we will notify users via their registered email address or through an in-app notice. Continued use of FIWB AI following notification constitutes acceptance of the updated policy.

Field	Details
Platform Name	FIWB AI
Operated By	SIOWA (Owais Sayyed & Siddhant Wagh)
Contact Email	fiwbai26@gmail.com
Policy Version	1.0
Effective Date	February 2026
Jurisdiction	India (subject to DPDP Act 2023 and applicable institutional regulations)

FIWB AI is committed to academic privacy and data minimization.

Your data exists to help you learn — for nothing else.